

Hashing

(N:12)

Open Addressing via Linear Probing: Insertion

- ▶ Compute the home bucket $L = h(K)$
- ▶ if $T[L]$ is not empty, consider the hash table as circular:

```
for( i = 0; i < m; i++ )  
    compute  $L = ( h(K) + i ) \% m$ ;  
    if  $T[L]$  is empty, put  $K$  there and stop
```

- ▶ If we can't find a position (the table is completely full), return an error message

An Example

- ▶ $D = 11$
- ▶ Add 58: Collision with 80
- ▶ Add 24
- ▶ Add 35: Collision with 24
- ▶ Insertion of 13 will join two clusters

			80				40			65
0	1	2	3	4	5	6	7	8	9	10

		24	80	58			40			65
0	1	2	3	4	5	6	7	8	9	10

		24	80	58	35		40			65
0	1	2	3	4	5	6	7	8	9	10

Searching for linear probing

- ▶ Search begins at the home bucket $h(k)$ for the key k
- ▶ Continues by examining successive buckets in the table by regarding the table as circular until one of the following happens
 1. A bucket containing the element with key k is reached; in this case, we found the element;
 2. An empty bucket is reached; in which case, the element is not found
 3. We return to the home bucket; in which case, the element is not found

Clustering in Linear Probing

- ▶ We call a block of contiguously occupied table entries a cluster
- ▶ Linear probing becomes slow when large clusters begin to form (primary clustering)
- ▶ For example:
 - ▶ Once $h(K)$ falls into a cluster, the cluster will definitely grow in size by 1
 - ▶ Larger clusters are easier targets for collision in the future
 - ▶ If two clusters are only separated by one entry, then inserting one key into a cluster can merge the two clusters

Deletion

- ▶ May require several movement: we cannot simply make the position empty. E.g., consider deleting 58

		24	80	58	35		40			65
0	1	2	3	4	5	6	7	8	9	10

- ▶ Move begins just after the bucket vacated by the deleted element. We need to *rehash* buckets one by one in the remainder of the cluster.
- ▶ Clearly, it may lead to a lot of movements involving at worst the whole table (of $O(m)$, where m is the size of the table)
- ▶ To reduce rehashing overhead, we can simply mark the entry “deleted,” which means treating it belonging to the cluster but skipping it in a cluster inspection
 - ▶ We need to distinguish it from the empty bucket at the cluster boundary

Introduction of a State Field in a Bucket

- ▶ A new state variable is inserted into each bucket



- ▶ The condition for unsuccessful search is that an *EMPTY* bucket is reached (not *DELETED* because it belongs to the cluster)
- ▶ After a while, almost all buckets have this status field set to *ACTIVE* or *DELETED*, and unsuccessful searches examine all buckets
- ▶ To improve performance, we must reorganize the table by, for example, rehashing into a new fresh table

Linear Probing Performance

- ▶ Let m be the number of buckets in the hash table
- ▶ n elements are present in the table
- ▶ Worst case search and insert time is $\Theta(n)$ (all n keys have the same home bucket)
- ▶ Average performance: U_n and S_n be the average number of buckets examined during an unsuccessful and successful search, respectively, and $\alpha = n/m$ be the load factor:

$$U_n \approx \frac{1}{2} \left[1 + \frac{1}{(1 - \alpha)^2} \right] \qquad S_n \approx \frac{1}{2} \left[1 + \frac{1}{1 - \alpha} \right]$$

An Application Example

- ▶ A hash table is to store up to 1,000 elements. Need to find its hash table size.
- ▶ Successful searches should require no more than 4 bucket examination on average and unsuccessful searches should examine no more than 50.5 buckets on average
 - ▶ i.e., $S_n \leq 4 \rightarrow \alpha \leq 6/7$
 - ▶ i.e., $U_n \leq 50.5 \rightarrow \alpha \leq 0.9$
- ▶ We hence require $\alpha = \min(6/7, 0.9) = 6/7$ and therefore $b \geq 1167$
- ▶ We choose D to be $37 \times 37 = 1369$ (no prime factors less than 20)

Quadratic Probing

- ▶ Insertion
- ▶ Compute $L = h(K)$
- ▶ Quadratic jump away from its home bucket
- ▶ If $T[L]$ is not empty:

```
for( i = 0; i < m; i++ )  
    compute  $L = ( h(K) + i*i ) \% m$ ;  
    if  $T[L]$  is empty, put K there and stop
```

- ▶ Helps to eliminate primary clustering
- ▶ However, if the table gets too full, this approach is not guaranteed of finding an empty slot!
- ▶ It may also never visit the home bucket again.

Double Hashing

- ▶ To alleviate the problem of primary clustering
- ▶ Use a second hash function h_2 when collision occurs
- ▶ Resolve collision by choosing the subsequent positions with a constant offset independent of the primary position
- ▶ Incrementally jump away from its home bucket in constant step size depending on the key
 - ▶ $H(K_i, 0) = h(K_i)$
 - ▶ $H(K_i, 1) = (H(K_i, 0) + h_2(K_i)) \bmod m$
 - ▶ $H(K_i, 2) = (H(K_i, 1) + h_2(K_i)) \bmod m$
 - ▶ ...
 - ▶ $H(K_i, m) = (H(K_i, m-1) + h_2(K_i)) \bmod m$

Choice of h_2

- ▶ For any key K , $h_2(K)$ must be relatively prime to the table size m
- ▶ Otherwise, we will only be able to examine a fraction of the table entries
- ▶ For example, if $h_2(K) = m/2$ (not relatively prime to m), then for $h(K) = 0$ we can only examine the entries $T[0]$ and $T[m/2]$ and nothing else!
- ▶ The only solution is to make m prime, and choose r to be a prime smaller than m , and set $h_2(K) = r - (K \bmod r)$
 - ▶ E.g., if $m = 37$, we may pick $r = 23$ and hence $h_2(K) = 23 - K \bmod 23$
 - ▶ We may as well use $r = 11$, and hence $h_2(K) = 11 - K \bmod 11$

Hash Performance

Strategies for improved performance:

1. Increase table capacity (less collisions)
2. Use a different hash function

► Hash table capacity

- Size of table is better to be at least 1.5 to 2 times the size of the number of items to be stored
- Otherwise probability of collisions would be very high

Comparison between BST and hash tables

BST	Hash tables
Comparison-based	Non-comparison-based
$O(\log n)$ time per operation	$O(1)$ time per operation
$O(n)$ space	$O(n)$ space
Keys stored in sorted order	Keys stored in arbitrary order
More operations are supported: min, max, neighbor, traversal	Only search, insert, delete
Can be augmented to support range queries	Do not support range queries
In C++: <code>std::map</code>	In C++: <code>std::unordered_map</code>

Q&A